

Titanic Logistic Regression basic

March 30, 2021

Logistic Regression on Titanic dataset

Sohini Mukherjee

30.03.2021

Importing Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[2]: train = pd.read_csv('E:/2.PYTHON-ML-BOOTCAMP/resources/13-Logistic-Regression/
↳titanic_train.csv')
```

```
[3]: train.head()
```

```
[3]: PassengerId  Survived  Pclass  \
0              1         0        3
1              2         1        1
2              3         1        3
3              4         1        1
4              5         0        3
```

```

                                Name      Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris   male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0    1
2                Heikkinen, Miss. Laina   female  26.0    0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)   female  35.0    1
4                Allen, Mr. William Henry   male  35.0    0
```

```

   Parch      Ticket    Fare Cabin Embarked
0      0   A/5 21171    7.2500   NaN        S
1      0   PC 17599   71.2833   C85        C
2      0  STON/O2. 3101282    7.9250   NaN        S
3      0    113803   53.1000  C123        S
4      0    373450    8.0500   NaN        S
```

Exploratory Data Analysis

Checking missing data

```
[4]: train.isnull()
```

```
[4]:
```


	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	\
0	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	False	
..	
886	False	False	False	False	False	False	False	False	False	
887	False	False	False	False	False	False	False	False	False	
888	False	False	False	False	False	True	False	False	False	
889	False	False	False	False	False	False	False	False	False	
890	False	False	False	False	False	False	False	False	False	

	Fare	Cabin	Embarked
0	False	True	False
1	False	False	False
2	False	True	False
3	False	False	False
4	False	True	False
..
886	False	True	False
887	False	False	False
888	False	True	False
889	False	False	False
890	False	True	False


```
[891 rows x 12 columns]
```


```
[5]: sns.heatmap(train.isnull(), cbar=False, yticklabels=False, cmap='viridis')
```

```
[5]: <AxesSubplot:>
```




```
[10]: #for classsification problems its a good idea to see the ratio of target labels.
      #checking who survived
      sns.set_style('whitegrid')
      sns.countplot(x='Survived', data=train, hue='Sex', palette='RdBu_r')
```

```
[10]: <AxesSubplot:xlabel='Survived', ylabel='count'>
```




```
[11]: sns.countplot(x='Survived', data= train, hue='Pclass')
```

```
[11]: <AxesSubplot:xlabel='Survived', ylabel='count'>
```




```
[13]: sns.displot(train['Age'].dropna(),bins=30)
```

```
[13]: <seaborn.axisgrid.FacetGrid at 0x2434881b3a0>
```




```
[14]: sns.countplot(x='SibSp', data=train)
```

```
[14]: <AxesSubplot:xlabel='SibSp', ylabel='count'>
```



```
[18]: train['Fare'].hist(bins=40,figsize=(10,5))
```

```
[18]: <AxesSubplot:>
```



```
[19]: import cufflinks as cf
      cf.go_offline()
```


```
[21]: train['Fare'].plot(kind='hist',bins=40)
```

Cleaning Data

From the heatmap it is evident that Age and Cabin columns have a lot of missing data. It is not wise to drop the Age column altogether. So we will fill in the missing data.

```
[24]: plt.figure(figsize=(10,5))
sns.boxplot(x='Pclass', y='Age', data=train)
```

```
[24]: <AxesSubplot:xlabel='Pclass', ylabel='Age'>
```



The people in the first class and second class are older than third class.

```
[30]: means = train.groupby('Pclass')['Age'].mean()
means
```

```
[30]: Pclass
1      38.233441
2      29.877630
3      25.140620
Name: Age, dtype: float64
```

```
[32]: def impute_age(cols):
    age = cols[0]
    pclass = cols[1]

    if pd.isnull(age):
```

```


    if pclass == 1:
        return 38
    elif pclass == 2:
        return 29
    else:
        return 25
else:
    return age

```

```
[33]: train['Age'] = train[['Age', 'Pclass']].apply(impute_age, axis=1)
```

```
[34]: sns.heatmap(train.isnull(), yticklabels=False, cbar=False)
```

```
[34]: <AxesSubplot:>
```



Cabin column has too many missing values so we are going to drop it.

```
[35]: train.drop('Cabin', axis= 1, inplace = True)
```

```
[36]: train.head()
```

```
[36]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	

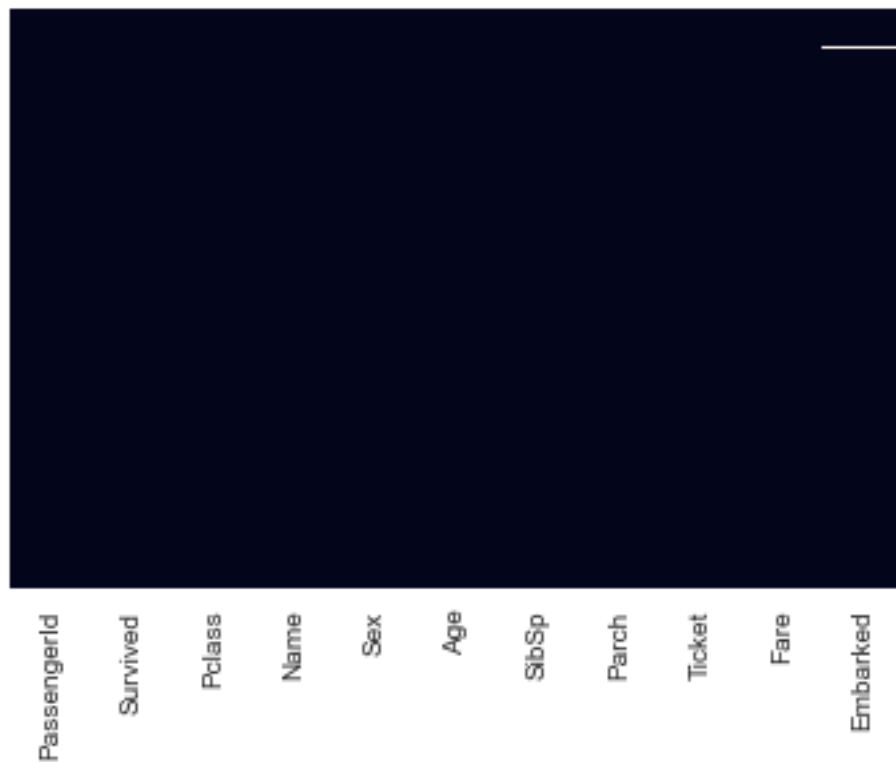
1	2	1	1
2	3	1	3
3	4	1	1
4	5	0	3

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Embarked
0	0	A/5 21171	7.2500	S
1	0	PC 17599	71.2833	C
2	0	STON/O2. 3101282	7.9250	S
3	0	113803	53.1000	S
4	0	373450	8.0500	S

```
[37]: sns.heatmap(train.isnull(), yticklabels=False, cbar=False)
```

```
[37]: <AxesSubplot:>
```

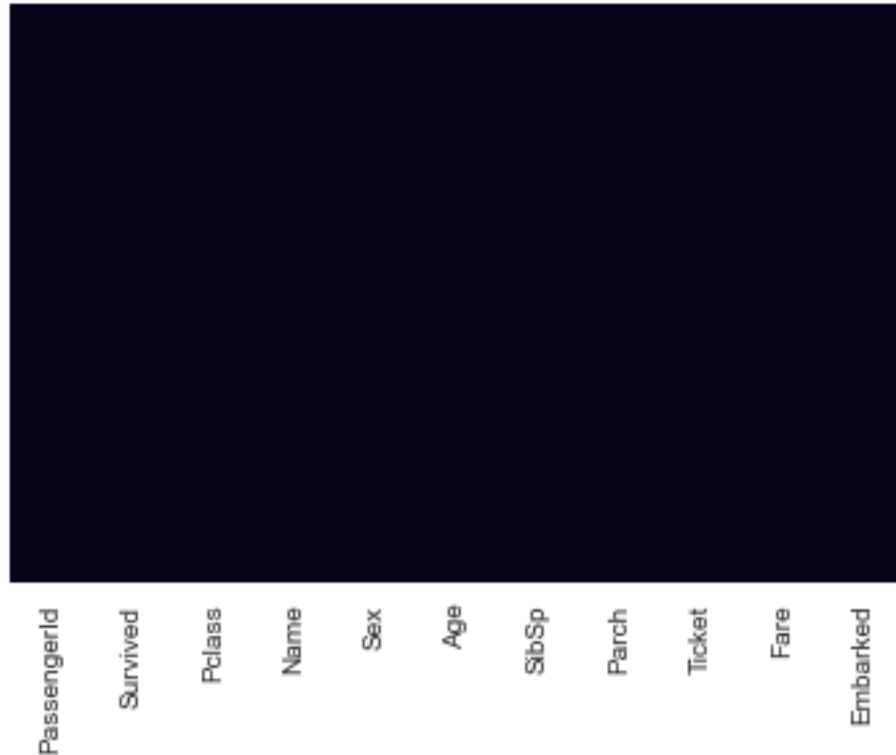


Dropping any remaining missing values.

```
[38]: train.dropna(inplace=True)
```

```
[39]: sns.heatmap(train.isnull(), yticklabels=False, cbar=False)
```

```
[39]: <AxesSubplot:>
```



Creating Dummy Variables for Sex and Embarked columns to apply Machine Learning

```
[42]: sex = pd.get_dummies(train['Sex'], drop_first=True) #using drop_first to avoid ↵  
      ↪multicollinearity problems
```

```
[43]: sex.head()
```

```
[43]:   male  
0     1  
1     0  
2     0  
3     0  
4     1
```

```
[44]: embark = pd.get_dummies(train['Embarked'], drop_first=True)
```

```
[45]: embark.head()
```

```
[45]:   Q  S
      0 0 1
      1 0 0
      2 0 1
      3 0 1
      4 0 1
```

```
[46]: train = pd.concat([train, sex, embark], axis=1)
```

```
[47]: train.head(1)
```

```
[47]:   PassengerId  Survived  Pclass   Name  Sex  Age  SibSp  \
      0         1         0        3  Braund, Mr. Owen Harris  male  22.0    1

      Parch  Ticket  Fare Embarked  male  Q  S
      0     0  A/5 21171   7.25         S    1 0 1
```

Dropping columns that are not usable

```
[51]: train.drop(['Sex', 'Embarked', 'Name', 'Ticket'], axis=1, inplace=True)
```

```
-----
KeyError                                Traceback (most recent call last)
<ipython-input-51-926d501f47cf> in <module>
----> 1 train.drop(['Sex', 'Embarked', 'Name', 'Ticket'], axis=1, inplace=True)

~\anaconda3\lib\site-packages\pandas\core\frame.py in drop(self, labels, axis,
↳ index, columns, level, inplace, errors)
   4161         weight 1.0    0.8
   4162         """
-> 4163         return super().drop(
   4164             labels=labels,
   4165             axis=axis,

~\anaconda3\lib\site-packages\pandas\core\generic.py in drop(self, labels, axis,
↳ index, columns, level, inplace, errors)
   3885         for axis, labels in axes.items():
   3886             if labels is not None:
-> 3887                 obj = obj._drop_axis(labels, axis, level=level,
↳ errors=errors)
   3888
   3889         if inplace:

~\anaconda3\lib\site-packages\pandas\core\generic.py in _drop_axis(self, labels,
↳ axis, level, errors)
```

```

3919             new_axis = axis.drop(labels, level=level, errors=errors)
3920         else:
-> 3921             new_axis = axis.drop(labels, errors=errors)
3922             result = self.reindex(**{axis_name: new_axis})
3923
~\anaconda3\lib\site-packages\pandas\core\indexes\base.py in drop(self, labels,
-> errors)
5280         if mask.any():
5281             if errors != "ignore":
-> 5282                 raise KeyError(f"{labels[mask]} not found in axis")
5283             indexer = indexer[~mask]
5284             return self.delete(indexer)

KeyError: "[ 'Sex' 'Embarked' 'Name' 'Ticket' ] not found in axis"

```

```
[55]: train.drop('PassengerId', axis=1, inplace=True)
```

```
[56]: train.head(1)
```

```
[56]:
```

	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	0	3	22.0	1	0	7.25	1	0	1

Machine Learning

```
[57]: X= train.drop('Survived', axis=1)
      y= train['Survived']
```

```
[58]: from sklearn.model_selection import train_test_split
```

```
[59]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
-> random_state=101)
```

```
[62]: from sklearn.linear_model import LogisticRegression
```

```
[66]: lg = LogisticRegression()
```

```
[67]: lg.fit(X_train, y_train)
```

C:\Users\ADMIN\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
[67]: LogisticRegression()
```

```
[68]: predictions = lg.predict(X_test)
```

```
[71]: from sklearn.metrics import classification_report
```

```
[72]: print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.83	0.91	0.87	163
1	0.84	0.70	0.76	104
accuracy			0.83	267
macro avg	0.83	0.81	0.82	267
weighted avg	0.83	0.83	0.83	267

```
[73]: from sklearn.metrics import confusion_matrix
```

```
[74]: confusion_matrix(y_test, predictions)
```

```
[74]: array([[149,  14],  
       [ 31,  73]], dtype=int64)
```

```
[ ]:
```